

CPE 453: Operating Systems
Assignment 4: Minix Secret Driver
Daniel Gutierrez (jguti151)

Executive Summary: In this assignment, I implemented a driver to interface with a Secret device. The Secret device can only hold one secret at a time.

Architecture

The driver implemented is intended to interact with the Secret device, `dev/Secret`. The Secret device can only hold one secret at a time. The Secret device is a character device.

Here is some nomenclature is used to help me better understand:

Empty - There is nothing in the buffer.

Full - The buffer is not empty.

Free - The secret is not owned by any user

Owned - The secret is owned by at least 1 user

The device's designed behavior is as follows:

Opening the Device:

If Free:

- Any process may open `/dev/Secret` for reading or writing (but not both).
 - Opening for read-write access will result in a permission denied error (EACCES).
- The owner of that process that opened it will then become the owner of the secret.
- Open for writing can only succeed if the secret is free.
 - This means it may only be opened for writing once.
- The secret held is of fixed size.
 - The default size is 8 KB, but can be reconfigured.
 - Attempts to write more into the device than will fit will result in an ENOSPC response.

If Owned:

- `/dev/Secret` may not be opened for writing once it is holding a secret.
 - Attempting to write to a full device results in an ENOSPC error.
- `/dev/Secret` may be opened for reading by a process owned by the owner of the secret.
 - Attempts to open a full secret for writing result in a device full error (ENOSPC).
 - Attempts to read a secret belonging to another user result in a permission denied error (EACCES).
- The secret persists as long as there are open file descriptors associated with it.

Closing the Device:

- The number of open file descriptors is decremented.
 - Once the last file descriptor is closed after a read, the device becomes empty and free.
 - The device is then available for new writes.

IOCTL Behavior:

- Changing ownership
 - The desired user to grant ownership to is determined and the new owner of the secret is set to the desired user.

Live Update behavior:

- The secret owner, the current buffer data, the writing and reading positions within the buffer, and the open counter are saved.
- These variables are then restored to the driver.

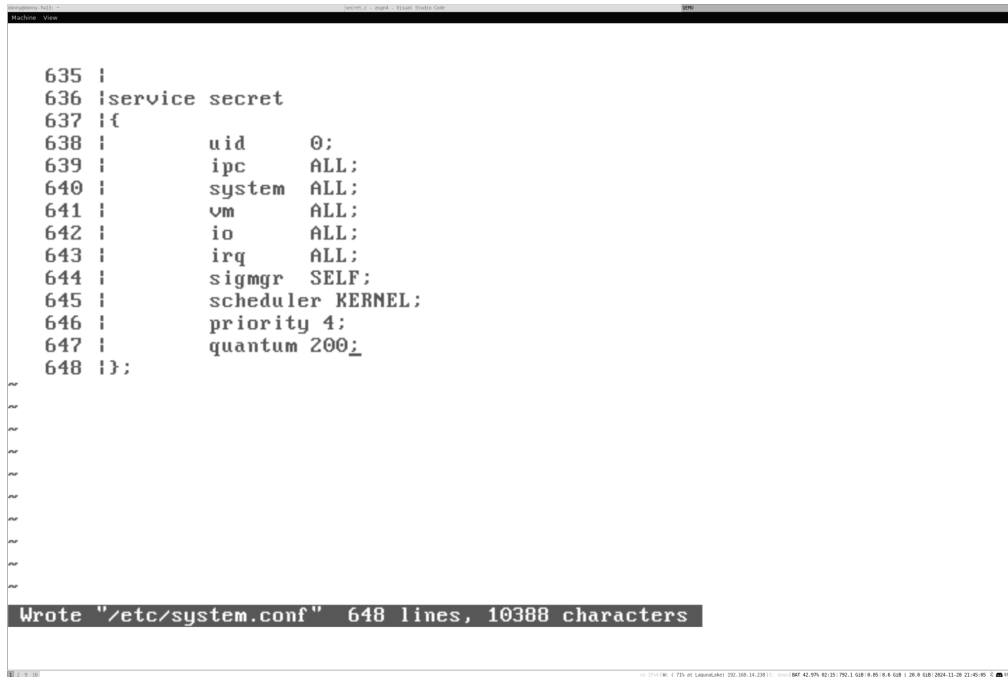
Driver Description

Dev Environment

I developed on MINIX 3.1.8 running on QEMU on a machine running Ubuntu 20.04.

Files Modified

/etc/system.conf



```
635 |
636 |service secret
637 |{
638 |    uid      0;
639 |    ipc      ALL;
640 |    system   ALL;
641 |    vm       ALL;
642 |    io       ALL;
643 |    irq      ALL;
644 |    sigmgr   SELF;
645 |    scheduler KERNEL;
646 |    priority 4;
647 |    quantum  200;
648 |};

Wrote "/etc/system.conf" 648 lines, 10388 characters
```

I added the above configuration for the system to recognize the service I was adding. I decided to make the service all powerful to make it as simple as possible. In hindsight, this is a very poor decision, but for the sake of the lab it works for what I need it to do even if it is a security risk. I chose a priority of 4 to designate it as a user-level process. The quantum was set to 200ms, based on information from [Individual Programming Assignment User Mode Scheduling in MINIX 3](#)^[1] indicating this as the default value.

```

4 | * This header file includes all other ioctl command code headers.
5 | */
6 |
7 | #ifndef _S_IOCTL_H
8 | #define _S_IOCTL_H
9 |
10 | /* A driver that uses ioctls claims a character for its series of commands
11 | * For instance: #define TCGETS _IOR('T', 8, struct termios)
12 | * This is a terminal ioctl that uses the character 'T'. The character(s)
13 | * used in each header file are shown in the comment following.
14 | */
15 |
16 | #include <sys/ioc_tty.h>          /* 'T' 't' 'k' */
17 | #include <net/ioctl.h>           /* 'n' */
18 | #include <sys/ioc_disk.h>        /* 'd' */
19 | #include <sys/ioc_file.h>        /* 'f' */
20 | #include <sys/ioc_memory.h>      /* 'm' */
21 | #include <sys/ioc_cmos.h>        /* 'c' */
22 | #include <sys/ioc_tape.h>        /* 'M' */
23 | #include <sys/ioc_scsi.h>        /* 'S' */
24 | #include <sys/ioc_sound.h>       /* 'S' */
25 | #include <sys/ioc_secret.h>      /* 'K' */
26 |
27 | #endif /* _S_IOCTL_H */

```

```
/usr/include/sys/ioc secret.h
```

```
1 #include <minix/ioctl.h>
2 #define SSGRANT_IOW('R', 1, uid_t)
3 #
```

I broke this assignment down into smaller parts to make it easier to complete: copying the hello example, modifying it to take in input, modifying it to hold the one secret, modifying it to be locked to 1 user, transferring users, then saving state.

Src Code

secret.c

```
#include "secret.h"

/* Function prototypes for secret driver */
PRIVATE char* secret_name(void);
PRIVATE int secret_open(struct driver* d, message* m);
PRIVATE int secret_close(struct driver* d, message* m);
PRIVATE struct device* secret_prepare(int device);
PRIVATE int secret_transfer(int procnr, int opcode,
                           u64_t position, iovec_t* iov, unsigned nr_req);
PRIVATE void secret_geometry(struct partition* entry);
PRIVATE int secret_ioctl(struct driver* d, message* m);

/* SEF functions */
PRIVATE void sef_local_startup(void);
PRIVATE int sef_cb_init(int type, sef_init_info_t* info);
PRIVATE int sef_cb_lu_state_save(int);
PRIVATE int lu_state_restore(void);

/* Entry points to the hello driver */
PRIVATE struct driver secret_tab = {
    secret_name,
    secret_open,
    secret_close,
    secret_ioctl,
    secret_prepare,
    secret_transfer,
    nop_cleanup,
    secret_geometry,
    nop_alarm,
    nop_cancel,
    nop_select,
    do_nop,
};

/* the device */
PRIVATE struct device secret_device;

/* state variables */
PRIVATE int fd_open_counter;
PRIVATE int is_readable;
PRIVATE int owned;
PRIVATE uid_t owner_uid;
PRIVATE int read_pos;
PRIVATE int write_pos;
PRIVATE char buffer[SECRET_SIZE];
```

```

/* gets the name of the device */
PRIVATE char* secret_name(void) {
    return "secret";
}

/* opens the device if free */
PRIVATE int secret_open(struct driver* d, message* m) {
    struct ucred caller_process;
    int reading, writing;
    /*
    In brightest code, in darkest compile, no relic of C89 shall beguile.
    Let those who cling to its ancient style, beware
    the future—it's worth the trial.

    ( courtesy of chatgpt for this poem but seriously i hate c89,
      c99 for life. let me declare variables anywhere)
    */

    getnucred(m->IO_ENDPT, &caller_process);
    reading = m->COUNT & R_BIT;
    writing = m->COUNT & W_BIT;

    if (!owned) {
        if (reading && writing) {
            printf("[OPEN] tring to RW while free\n");
            return EACCES;
        }
        else if (reading) {
            is_readable = FALSE;
        }
        else { /* writing to it so it becomes readable" */
            is_readable = TRUE;
        }
        owner_uid = caller_process.uid;
        owned = TRUE;
    }
    else {
        if (writing) { /* can only write once */
            printf("[OPEN] trying to write while owned\n");
            return EACCES;
        }
        else {
            if (owner_uid != caller_process.uid) {
                printf("[OPEN] trying to read while owned by someone else\n");
                return EACCES;
            }
            is_readable = FALSE;
        }
    }
}

```

```

    }

    fd_open_counter++;
    return OK;
}

/* closes the device */
PRIVATE int secret_close(struct driver* d, message* m) {
    struct ucred caller_process;

    getnucred(m->IO_ENDPT, &caller_process);

    fd_open_counter--;

    if (fd_open_counter == 0 && !is_readable) {
        owner_uid = NO_OWNER_ID;
        owned = FALSE;
        memset(buffer, 0, SECRET_SIZE);
        write_pos = 0;
        read_pos = 0;
    }

    return OK;
}

/* switch owners */
PRIVATE int secret_ioctl(struct driver* d, message* m) {
    uid_t new_owner_uid;
    int ret;

    if (m->REQUEST != SSGRANT) {
        return ENOTTY;
    }

    ret = sys_safecopyfrom(m->IO_ENDPT, (vir_bytes)m->IO_GRANT, 0,
                          (vir_bytes)&new_owner_uid, sizeof(uid_t), D);
    owner_uid = new_owner_uid;

    return ret;
}

/* prepare the device (stolen from hello) */
PRIVATE struct device* secret_prepare(int device) {
    secret_device.dv_base.lo = 0;
    secret_device.dv_base.hi = 0;
    secret_device.dv_size.lo = 0;
    secret_device.dv_size.hi = 0;

```



```

    return &secret_device;
}

/* transfer data to/from the device */
PRIVATE int secret_transfer(int procnr, int opcode,
                           u64_t position, iovec_t* iov, unsigned nr_req) {
    int bytes;
    int ret;

    switch (opcode) {
        case DEV_GATHER_S: /* read from device */
            bytes = write_pos - read_pos < iov->iov_size ?
                    write_pos - read_pos : iov->iov_size;
            if (bytes <= 0) { /* unsuccessful read but continue on */
                return OK;
            }

            if (ret = sys_safecopyto(procnr, iov->iov_addr, 0,
                                    (vir_bytes)(buffer + read_pos),
                                    bytes, D) != OK) {
                return ret;
            }

            iov->iov_size -= bytes;
            read_pos += bytes;
            break;

        case DEV_SCATTER_S: /* write to device */
            bytes = SECRET_SIZE - write_pos < iov->iov_size ?
                    SECRET_SIZE - write_pos : iov->iov_size;
            if (bytes <= 0) { /* unsuccessful write but continue on */
                return ENOSPC;
            }

            if (ret = sys_safecopyfrom(procnr, iov->iov_addr, 0,
                                      (vir_bytes)(buffer + write_pos),
                                      bytes, D) != OK) {
                return ret;
            }

            iov->iov_size -= bytes;
            write_pos += bytes;
            break;

        default:
            return EINVAL;
    }
}

```

```

    return ret;
}

/* stolen from hello */
PRIVATE void secret_geometry(struct partition* entry) {
    printf("hello_geometry()\n");
    entry->cylinders = 0;
    entry->heads = 0;
    entry->sectors = 0;

    return;
}

/* save all the variables to be restored later */
PRIVATE int sef_cb_lu_state_save(int state) {
    SecretState_t cur_secret;
    cur_secret.fd_open_counter = fd_open_counter;
    cur_secret.owner_uid = owner_uid;
    cur_secret.owned = owned;
    cur_secret.write_pos = write_pos;
    cur_secret.read_pos = read_pos;
    memcpy(cur_secret.buffer, buffer, SECRET_SIZE);

    ds_publish_mem(SECRET_STATE_NAME, (char*)&cur_secret,
        sizeof(SecretState_t), DSF_OVERWRITE);

    return OK;
}

PRIVATE int lu_state_restore() {
    SecretState_t restored_secret;
    size_t len;

    ds_retrieve_mem(SECRET_STATE_NAME, (char*)&restored_secret, &len);

    fd_open_counter = restored_secret.fd_open_counter;
    owner_uid = restored_secret.owner_uid;
    owned = restored_secret.owned;
    write_pos = restored_secret.write_pos;
    read_pos = restored_secret.read_pos;
    memcpy(buffer, restored_secret.buffer, SECRET_SIZE);

    ds_delete_mem(SECRET_STATE_NAME);

    return OK;
}

```

```

/* stolen from hello */
PRIVATE void sef_local_startup() {
    /* Register init callbacks. Use the same function for all event types */
    sef_setcb_init_fresh(sef_cb_init);
    sef_setcb_init_lu(sef_cb_init);
    sef_setcb_init_restart(sef_cb_init);

    /* Register live update callbacks */
    /* - Agree to update immediately when LU is requested in a valid state. */
    sef_setcb_lu_prepare(sef_cb_lu_prepare_always_ready);
    /* - Support live update starting from any standard state. */
    sef_setcb_lu_state_isvalid(sef_cb_lu_state_isvalid_standard);
    /* - Register a custom routine to save the state. */
    sef_setcb_lu_state_save(sef_cb_lu_state_save);
    /* Let SEF perform startup. */
    sef_startup();

    return;
}

/* initialize the driver */
PRIVATE int sef_cb_init(int type, sef_init_info_t *info) {
    int do_announce_driver = TRUE;

    switch(type) {
        case SEF_INIT_FRESH:
            fd_open_counter = 0;
            is_readable = FALSE;
            owned = FALSE;
            owner_uid = NO_OWNER_ID;
            read_pos = 0;
            write_pos = 0;
            memset(buffer, 0, SECRET_SIZE);
            break;
        case SEF_INIT_LU:
            lu_state_restore();
            do_announce_driver = FALSE;
            printf("Secret driver: live update\n");
            break;
        case SEF_INIT_RESTART:
            printf("Secret driver: restart\n");
            break;
    }

    /* announce when up */
    if (do_announce_driver) {

```

```
        driver_announce();
    }
    /* been initialized properly */
    return OK;
}

/* stolen from hello */
PUBLIC int main(int argc, char **argv) {
    /*
     * Perform initialization.
     */
    sef_local_startup();

    /*
     * Run the main loop.
     */
    driver_task(&secret_tab, DRIVER_STD);

    return OK;
}
```

Secret.h

```
#ifndef SECRET_H
#define SECRET_H

#include <minix/drivers.h>
#include <minix/driver.h>
#include <stdio.h>
#include <stdlib.h>
#include <minix/ds.h>
#include <sys/ioc_secret.h>

#define SECRET_MSG ("I am a secret hehe")
#define NO_OWNER_ID  (-1)

#ifndef SECRET_SIZE
#define SECRET_SIZE (8192)
#endif

/* struct for simplifying saving state */
typedef struct SecretState_t {
    int fd_open_counter;
    int is_readable;
    int owned;
    uid_t owner_uid;
    int read_pos;
    int write_pos;
    char buffer[SECRET_SIZE];
} SecretState_t;
#define SECRET_STATE_NAME ("secret_state")

#endif /* secret.h */
```

Running Behavior

```
# whoami
root
# service up $PWD/secret -dev /dev/Secret
# cat /dev/Secret
# echo "this is a secret" > /dev/Secret
# echo "this should fail" > /dev/Secret
cannot create /dev/Secret: No space left on device
# cat /dev/Secret
this is a secret
# cat /dev/Secret
# echo "tis but another secret" > /dev/Secret
# su danny
$ cat /dev/Secret
cat: /dev/Secret: Permission denied
$ cat > /dev/Secret
cannot create /dev/Secret: No space left on device
$ exit
# cat /dev/Secret
tis but another secret
# su danny
$ echo "danny's secret" > /dev/Secret
$ exit
# cat /dev/Secret
cat: /dev/Secret: Permission denied
#
```

```
# cat /dev/Secret
tis but another secret
# su danny
$ echo "danny's secret" > /dev/Secret
$ exit
# cat /dev/Secret
cat: /dev/Secret: Permission denied
# su danny
$ cat /dev/Secret
danny's secret
$ exit
# cd /usr/src/drivers/secrets/
# ls -l
total 85
-rw-r--r-- 1 root operator 1892 Nov 26 12:09 .depend
-rw-r--r-- 1 root operator 166 Nov 26 11:20 Makefile
-rw-r--r-- 1 root operator 16384 Nov 28 04:14 bigfile.txt
-rwxr-xr-x 1 root operator 52152 Nov 26 12:09 secret
-rw-r--r-- 1 root operator 8158 Nov 26 12:09 secret.c
-rw-r--r-- 1 root operator 1892 Nov 26 12:09 secret.d
-rw-r--r-- 1 root operator 606 Nov 26 11:20 secret.h
-rw-r--r-- 1 root operator 3913 Nov 26 12:09 secret.o
# cat bigfile.txt > /dev/Secret
cat: standard output: No space left on device
#
```

The devicer is working as intended verified by running the same run as in the assignment spec. As I got the same result as in the sample run and passed all the test on test suite I deemed the driver as working.

Problems Encountered

The only problem I encountered was where to place the ioctl include for ioc_secret. I initially edited /usr/src/include/sys/ioctl.h and when trying to compile my driver I encountered error: unable to open ioc_secret.h. After running find / -name "ioctl.h" and placing it in usr/src/include/sys/ioctl.h it worked.